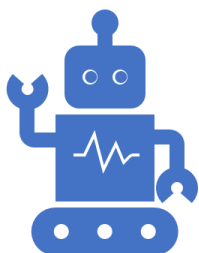




UiO : **University of Oslo**

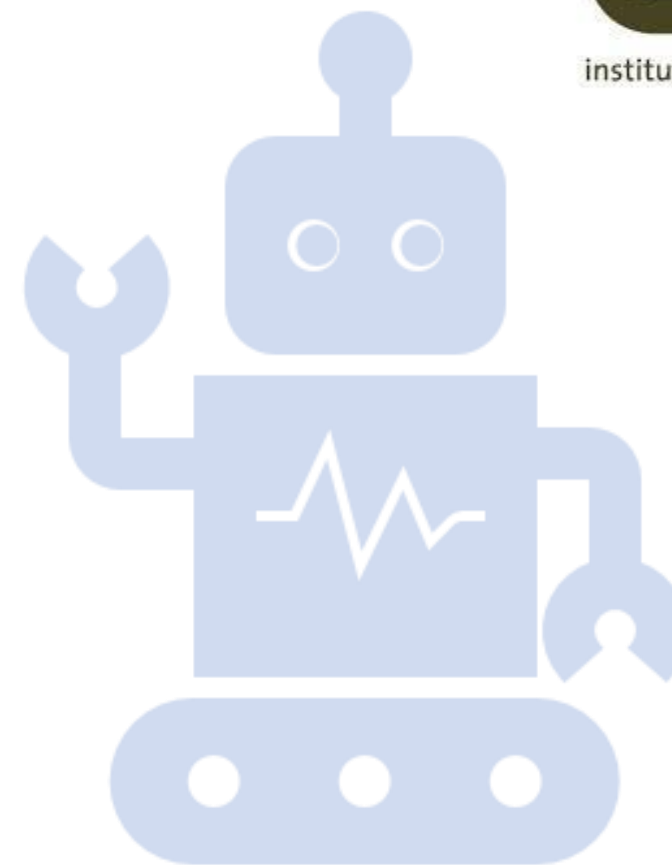


IN3050/IN4050 - Introduction to Artificial Intelligence and Machine Learning

Lecture 8


Multi-layer neural networks and backpropagation

Jan Tore Lønning



8.1 Feed-forward Neural networks

IN3050/IN4050 Introduction to Artificial Intelligence
and Machine Learning



Today


1. Feed-forward neural networks (Multi-layer Perceptron)
2. Matrix representations of neural networks
3. The Backpropagation Algorithm
4. Finer details
5. More on Evaluation

The neural inspiration

- So far inspired by one neuron
- That does not make intelligence

The human brain, roughly


- 10^{11} Neurons
- 10^{14} Synapses
- The strength is the interactions
- Neural Networks



<https://simple.wikipedia.org/wiki/Neuron#/media/File:Neuron.svg>


Artificial Neural Networks

- Inspired by the brain
- Does not pretend to be a model of the brain
- The simplest model is the
 - Feed forward network, also called
 - Multi-layer Perceptron




Feed forward network

- An input layer
- An output layer: the predictions
- One or more hidden layers
- Connections from nodes in one layer to nodes in the next layer (from left to right)
- The connections are marked with weights



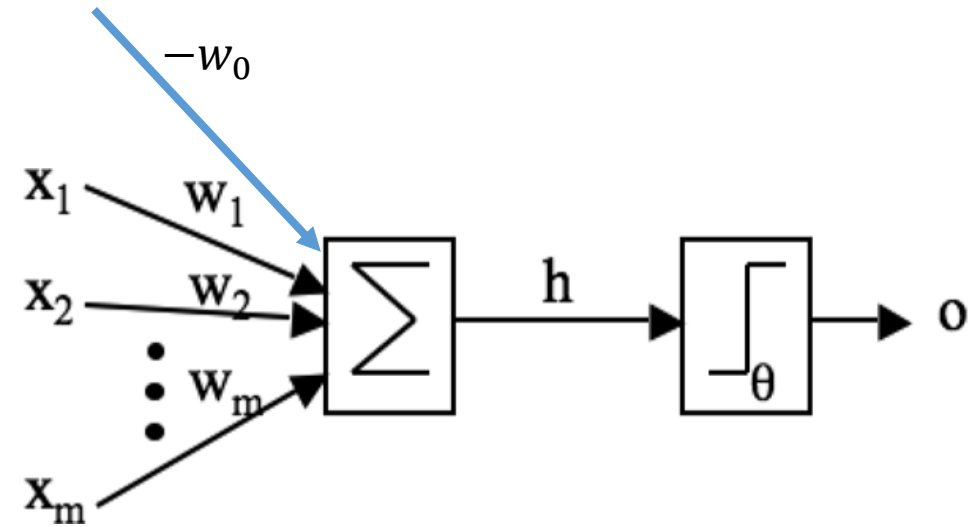
Going forwards (predictions)

- There is one input node for each feature/dimension in an input vector: (x_1, x_2, \dots, x_m)
- In addition, an input bias node $x_0 = -1$
- The input values are multiplied with the weights and summed into each hidden node.
- There is some processing in the hidden node.
- The output values of the hidden nodes are fed to the next layer.
- (etc.)



One hidden unit


1. First sum of weighted inputs :
 - $z = \sum_{i=0}^m w_i x_i = \mathbf{w} \cdot \mathbf{x}$
 2. Then the result is run through an activation function, g to produce $g(z) = g(\mathbf{w} \cdot \mathbf{x})$
- The activation function could be the step function,
 - c.f. the XOR-example:
 - Marsland sec 3.4.2 & start of ch. 4



It is the non-linearity of the activation function which makes it possible for MLP to predict non-linear decision boundaries

A differentiable activation function

- It is unclear how to update the weights if g isn't differentiable
- One option is to use the logistic (sigmoid) function
 - $y = \sigma(z) = \frac{1}{1+e^{-\vec{w} \cdot \vec{x}}}$
 - Differentiable
 - $y' = y(1 - y)$
- (There are alternative activation functions.)




One hidden node

1. First sum of weighted inputs:

- $z = \sum_{i=0}^m w_i x_i = \mathbf{w} \cdot \mathbf{x}$


2. Then

- $y = g(z) = \sigma(z) = \frac{1}{1+e^{-\vec{w} \cdot \vec{x}}}$




Going forwards (predictions)

- After the processing in the hidden layer, the output is taken as input to the next layer
- One must also add a bias term at this layer.
 - Observe that this has to be done:
 - During processing
 - E.g., over again each time we process the same training item




Output layer

- Several possibilities, depending on the task, including:
 - Regression
 - Binary classification
 - Multi-label classification
 - Multi-class classification
- From the last layer to the output layer is like the same tasks without multiple layers!
- c.f. Marsland, sec. 4.2.3




1. Regression

- One output node
- No activation function in the output layer
 - = activation function is the identity function
- Observe that this can predict non-linear functions!




2. Binary classification

- One output node
- Logistic activation function in the output layer
- Similar to logistic regression
- Can produce non-linear decision boundaries




3. Multi-label classification


- Several output nodes
- Logistic activation function
- Can be made multi-class classification by one vs. rest.
- The model Marsland considers



4. Multi-class classification

- Several output nodes
- Sum the weighted inputs at each nodes
- The sums are brought together in the soft-max






8.2 Matrix representations

IN3050/IN4050 Introduction to Artificial Intelligence
and Machine Learning

Representing the connections


$$\begin{bmatrix} x_0 & \textcircled{x_1} & x_2 & \cdots & x_m \end{bmatrix} \begin{bmatrix} w_{0,1} & w_{0,2} & \cdots & w_{0,n} \\ w_{1,1} & \textcircled{w_{1,2}} & \cdots & w_{1,n} \\ w_{2,1} & w_{2,2} & \cdots & w_{2,n} \\ \vdots & \vdots & \ddots & \vdots \\ w_{m,1} & w_{m,2} & \cdots & w_{m,n} \end{bmatrix} = \begin{bmatrix} z_1 & \textcircled{z_2} & \cdots & z_n \end{bmatrix}$$

- We use a matrix to represent the connections
- Element $w_{i,j}$ is the connection:
 - from node i
 - to node j
- (Beware, some texts do it differently)




Connections going into a node

$$\begin{bmatrix} x_0 & x_1 & x_2 & \cdots & x_m \end{bmatrix} \begin{bmatrix} w_{0,1} & w_{0,2} & \cdots & w_{0,n} \\ w_{1,1} & w_{1,2} & \cdots & w_{1,n} \\ w_{2,1} & w_{2,2} & \cdots & w_{2,n} \\ \vdots & \vdots & \ddots & \vdots \\ w_{m,1} & w_{m,2} & \cdots & w_{m,n} \end{bmatrix} = \begin{bmatrix} z_1 & z_2 & \cdots & z_n \end{bmatrix}$$



Connections going out of a node

$$\begin{bmatrix} x_0 & \boxed{x_1} & x_2 & \cdots & x_m \end{bmatrix} \begin{bmatrix} w_{0,1} & w_{0,2} & \cdots & w_{0,n} \\ \boxed{w_{1,1} & w_{1,2} & \cdots & w_{1,n}} \\ w_{2,1} & w_{2,2} & \cdots & w_{2,n} \\ \vdots & \vdots & \ddots & \vdots \\ w_{m,1} & w_{m,2} & \cdots & w_{m,n} \end{bmatrix} = \begin{bmatrix} z_1 & z_2 & \cdots & z_n \end{bmatrix}$$



Batch-processing

$$\begin{bmatrix} x_{1,0} & x_{1,1} & x_{1,2} & \cdots & x_{1,m} \\ x_{2,0} & x_{2,1} & x_{2,2} & \cdots & x_{2,m} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ x_{N,0} & x_{N,1} & x_{N,2} & \cdots & x_{N,m} \end{bmatrix} \begin{bmatrix} w_{0,1} & w_{0,2} & \cdots & w_{0,n} \\ w_{1,1} & w_{1,2} & \cdots & w_{1,n} \\ w_{2,1} & w_{2,2} & \cdots & w_{2,n} \\ \vdots & \vdots & \ddots & \vdots \\ w_{m,1} & w_{m,2} & \cdots & w_{m,n} \end{bmatrix} = \begin{bmatrix} z_{1,1} & z_{1,2} & \cdots & z_{1,n} \\ z_{2,1} & z_{2,2} & \cdots & z_{2,n} \\ \vdots & \vdots & \ddots & \vdots \\ z_{N,1} & z_{N,2} & \cdots & z_{N,n} \end{bmatrix}$$

- In batch-processing we can multiply by weights and (i) sum the results for (iii) each input item, and (ii) each hidden node in one operation
- Three nested loops by just: XW

Activation function

$$\begin{bmatrix} x_{1,0} & x_{1,1} & x_{1,2} & \cdots & x_{1,m} \\ x_{2,0} & x_{2,1} & x_{2,2} & \cdots & x_{2,m} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ x_{N,0} & x_{N,1} & x_{N,2} & \cdots & x_{N,m} \end{bmatrix} \begin{bmatrix} w_{0,1} & w_{0,2} & \cdots & w_{0,n} \\ w_{1,1} & w_{1,2} & \cdots & w_{1,n} \\ w_{2,1} & w_{2,2} & \cdots & w_{2,n} \\ \vdots & \vdots & \ddots & \vdots \\ w_{m,1} & w_{m,2} & \cdots & w_{m,n} \end{bmatrix} = \begin{bmatrix} z_{1,1} & z_{1,2} & \cdots & z_{1,n} \\ z_{2,1} & z_{2,2} & \cdots & z_{2,n} \\ \vdots & \vdots & \ddots & \vdots \\ z_{N,1} & z_{N,2} & \cdots & z_{N,n} \end{bmatrix}$$

- Each $z_{i,j}$ is passed through the activation function: $y_{i,j} = g(z_{i,j})$
- In NumPy this can be done by one operation: $g(XW)$
- Reminder: g may be the logistic function, but doesn't have to
 - i.e., $g(z_{i,j}) = \sigma(z_{i,j}) = \frac{1}{1+e^{-z_{i,j}}}$

Footnote: Notation


- Half of all texts follow us and Marsland with respect to notation
- The other half does differently

	We	Them
Connection from node i to node j	$w_{i,j}$	$w_{j,i}$
Data and weights	XW	WX
Applying activation function	$g(XW)$	$g(WX)$

- It amounts to the same.
- But don't mix them up!

8.3 Learning by Back-propagation

IN3050/IN4050 Introduction to Artificial Intelligence
and Machine Learning



Background

Marsland (p.74), “...just three things that you need to know...”:

1. If $f(x) = \frac{1}{2}x^2$ then $f'(x) = x$
2. If $f(x) = c$ then $f'(x) = 0$
3. If $f(x) = h(g(x))$ then $f'(x) = h'(g(x))g'(x)$ (the chain rule)

He forgot

4. If $y = \sigma(z) = \frac{1}{1+e^{-\vec{w} \cdot \vec{x}}}$, then $y' = y(1 - y)$


In addition

We will make use of the following which we have already seen:

- The logistic regression model
- Gradient descent
- GD applied to
 - Linear regression
 - Logistic regression
- Loss-functions:
 - MSE, Cross-Entropy

Training

- Given a set of training instances
 - $\{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_N, t_N)\}$:
- Forwards:
 - Run them forwards and get predictions
 - $\{y_1, y_2, \dots, y_N\}$
- Backwards
 - Use a suitable loss function and compare these to the target values
 - $\{t_1, t_2, \dots, t_N\}$
 - Apply gradient descent to update the weights (partial derivatives)



How do we update the weights

Last layer


- (easy)
- Like the same problems for linear regression or logistic regression without a hidden layer

The first layer

- The big question:
- How do we update the first layer?
- We don't have a loss (error) here

Solution: Backpropagation

- Let's be a little more formal
- Let the matrix **V** be the connections from *input* to *hidden* and **W** from *hidden* to *output*
 - $\dim(V) = ((m + 1) \times k)$
 - $\dim(W) = ((k + 1) \times n)$
- Activation functions:
 - Hidden layers: g
 - Output layer: f




- Let us in the following consider SGD where we update for one input $\mathbf{x} = (x_1, x_2, \dots, x_m)$

Forwards (notation)

- Add bias and send
 - $\mathbf{x}^+ = (\mathbf{x}_0, x_1, \dots, x_m)$
 - through the first layer to get
 - $\mathbf{h} = \mathbf{x}^+ \mathbf{V} = (h_1, h_2, \dots, h_k)$, where
 - $h_j = \sum_{i=0}^m x_i v_{i,j}$
 - k is the number of hidden nodes
 - Apply activation function to get
 - $\mathbf{a} = g(\mathbf{h}) = (a_1, a_2, \dots, a_k)$,
 - where $a_j = g(h_j)$
- Add bias and send
 - $\mathbf{a}^+ = (\mathbf{a}_0, a_1, a_2, \dots, a_k)$
 - through the second layer to get
 - $\mathbf{z} = \mathbf{a}^+ \mathbf{W} = (z_1, z_2, \dots, z_n)$, where
 - $z_j = \sum_{i=0}^k a_i w_{i,j}$
 - n is the number of output nodes
 - Apply activation function to get
 - $\mathbf{y} = f(\mathbf{z}) = (y_1, y_2, \dots, y_n)$,
 - where $y_j = f(z_j)$

Backwards: 1. Regression

- We will consider various output tasks, starting with the simple regression
- There is only one output node
- The output activation function, f , is identity



Backwards: Update last layer

- For loss, we use MSE, or , as Marsland, the simpler

Sum of Squares Error (SE): $L_{SE}(\mathbf{t}, \mathbf{y}) = \frac{1}{2} \sum_{j=1}^N (t_j - y_j)^2$

- (The index j here, runs over the input items. There is only one output node)
- We have seen that
- $\frac{\partial}{\partial w_{i,1}} L_{SE}(\mathbf{t}, \mathbf{y}) = \frac{\partial}{\partial \mathbf{y}} L_{SE}(\mathbf{t}, \mathbf{y}) \left(\frac{\partial}{\partial w_{i,1}} \mathbf{y} \right) = \sum_{j=1}^N ((t_j - y_j)(-a_{j,i}))$
- For SGD where we update for one input $\mathbf{x} = (x_1, x_2, \dots, x_m)$
- $\frac{\partial}{\partial w_{i,1}} L_{SE}(\mathbf{t}, \mathbf{y}) = \frac{\partial}{\partial \mathbf{y}} L_{SE}(\mathbf{t}, \mathbf{y}) \left(\frac{\partial}{\partial w_{i,1}} \mathbf{y} \right) = (t - y)(-a_i) = (y - t)(a_i)$

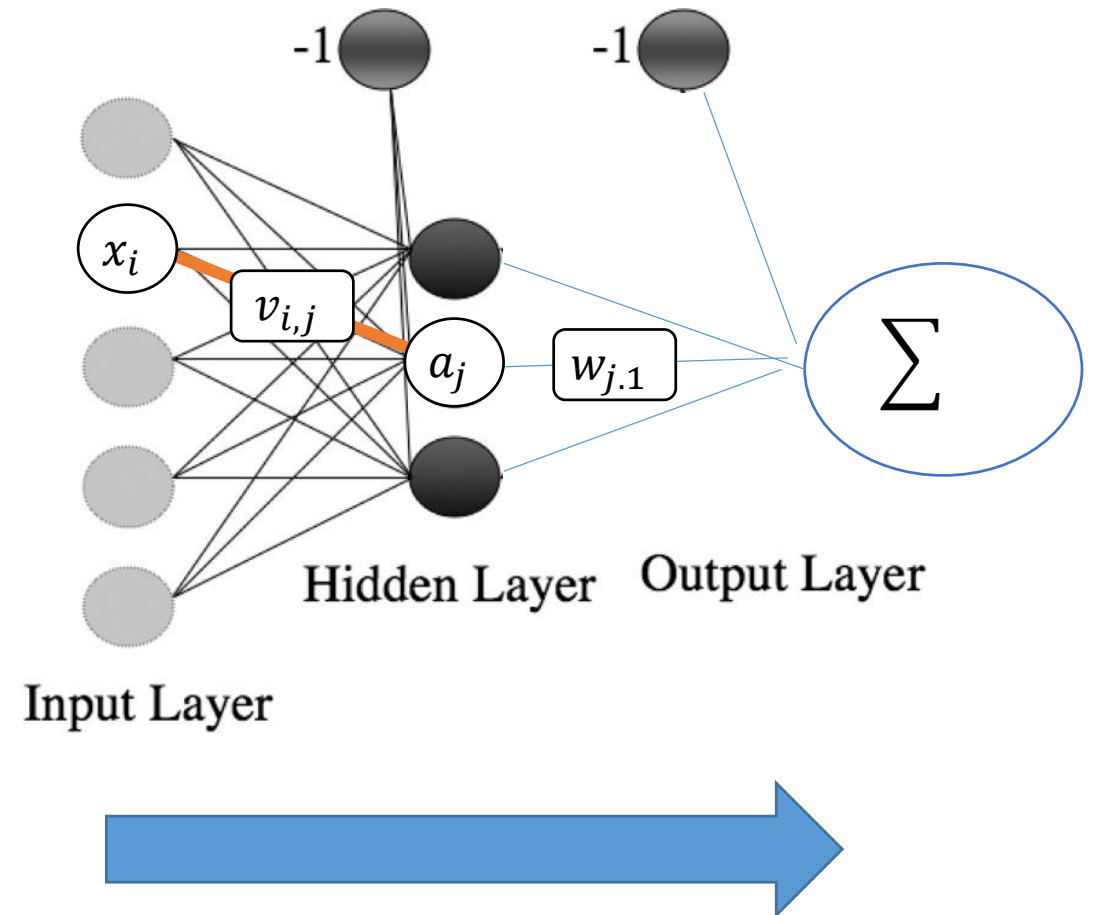
Backwards: Update last layer, ctd.

- $\frac{\partial}{\partial w_{i,1}} L_{SE}(t, y) = (y - t)(a_i)$
- We know from lect. 6 how to update this (\mathbf{a} corresponds to \mathbf{x} then)
- But wait!
- We first have to find how to update the first layer.



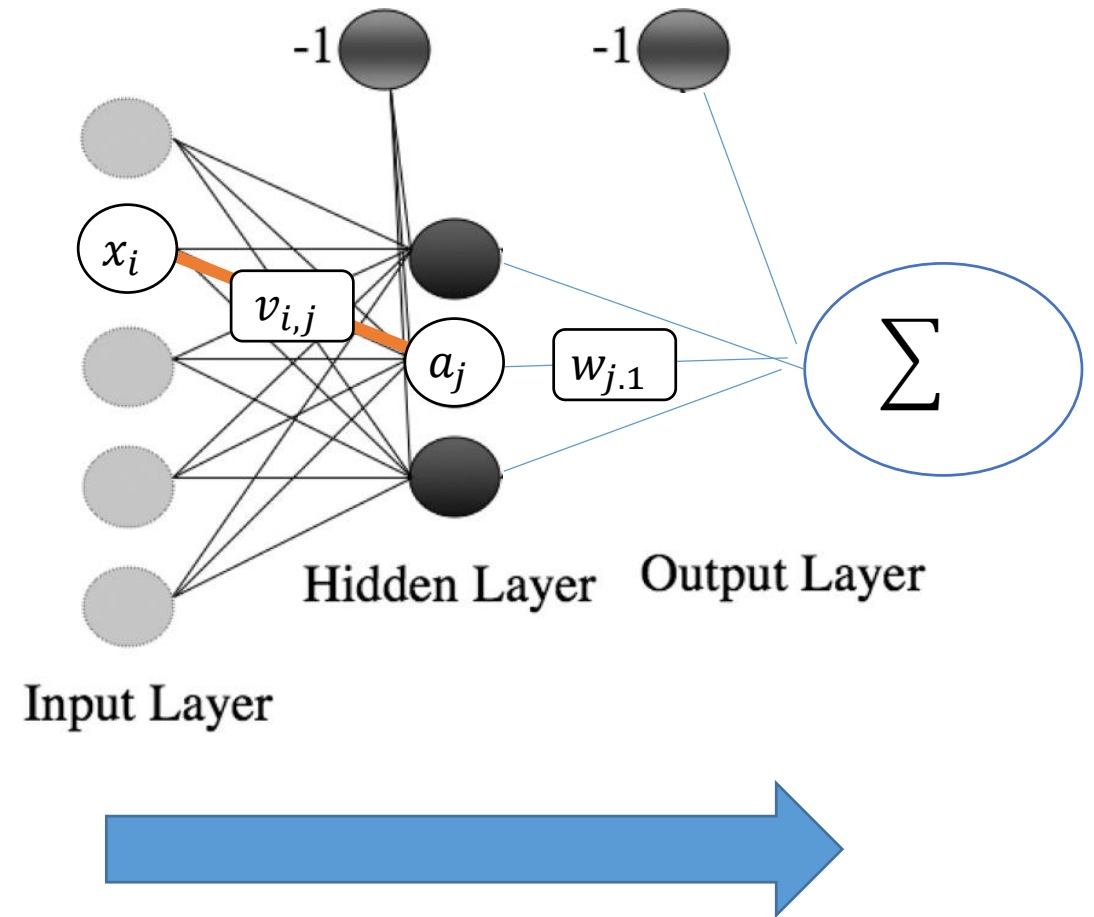
Backwards: Update first layer: V , 1

- $\mathbf{y} = f(\mathbf{z}) = \mathbf{z}$, where $\mathbf{z} = \mathbf{a}^+ \mathbf{W}$
- $\mathbf{a} = g(\mathbf{h})$, where $\mathbf{h} = \mathbf{x}^+ \mathbf{V}$
- $\frac{\partial}{\partial v_{i,j}} L_{SE}(\mathbf{t}, \mathbf{y}) =$
- $\frac{\partial}{\partial \mathbf{a}} L_{SE}(\mathbf{t}, \mathbf{y}) \left(\frac{\partial}{\partial v_{i,j}} \mathbf{a} \right) =$
- $\frac{\partial}{\partial a_j} L_{SE}(\mathbf{t}, \mathbf{y}) \left(\frac{\partial}{\partial v_{i,j}} a_j \right)$
- because $\left(\frac{\partial}{\partial v_{i,j}} a_k \right) = \mathbf{0}$ for $k \neq j$




Backwards: Update first layer: V, 2

- $\mathbf{y} = f(\mathbf{z}) = \mathbf{z}$, where $\mathbf{z} = \mathbf{a}^+ \mathbf{W}$
- $\frac{\partial}{\partial a_j} L_{SE}(t, y) = \frac{\partial}{\partial y} L_{SE}(t, y) \left(\frac{\partial}{\partial a_j} y \right) = (t - y)(-w_{j,1}) = (y - t)(\mathbf{w}_{j,1})$
- Observe similarities and differences to
- $\frac{\partial}{\partial w_{i,1}} L_{SE}(t, y) = (y - t)(\mathbf{a}_i)$
- We call the common part: $(y - t)$ for the **delta term** $\delta_o(\kappa)$ of the end node κ .



Backwards: Update first layer: V, 3

- $\mathbf{a} = g(\mathbf{h})$, where $\mathbf{h} = \mathbf{x}^+ \mathbf{V}$
- $\left(\frac{\partial}{\partial v_{i,j}} a_j \right) = \left(\frac{\partial}{\partial \mathbf{h}} g \right) \left(\frac{\partial}{\partial v_{i,j}} \mathbf{h} \right) = \left(\frac{\partial}{\partial h_j} g \right) \left(\frac{\partial}{\partial v_{i,j}} h_j \right)$
- $\frac{\partial}{\partial v_{i,j}} h_j = x_i$
- If $a_j = g(h_j) = \sigma(h_j)$, then
 - $\left(\frac{\partial}{\partial h_j} g \right) = a_j(1 - a_j)$
 - $\left(\frac{\partial}{\partial v_{i,j}} a_j \right) = a_j(1 - a_j)x_i$




Backwards: Update first layer: V , 4

- $\mathbf{y} = f(\mathbf{z}) = \mathbf{z}$, where $\mathbf{z} = \mathbf{a}^+ \mathbf{W}$
- $\mathbf{a} = g(\mathbf{h})$, where $\mathbf{h} = \mathbf{x}^+ \mathbf{V}$
- $\frac{\partial}{\partial v_{i,j}} L_{SE}(t, y) = \frac{\partial}{\partial a_j} L_{SE}(t, y) \left(\frac{\partial}{\partial v_{i,j}} a_j \right) =$

$\delta_o(\kappa)$

- $(y - t)(w_{j,1})a_j(1 - a_j)x_i$

δ -term at the node
marked with a_j




Putting it together: the Algorithm

- Use the loss function and the derivative of the activation function to compute the delta term at the final node(s), here: $\delta_o(\kappa_1) = (y - t)$
- Compute the delta terms for each node in the hidden layer, from the delta term(s) and the hidden layer and the weights at the connections
 - here: $\delta(hidden_j) = \delta_o(\kappa_1) (w_{j,1}) a_j (1 - a_j)$
- Update the weights by the deltas:
 - $w_{i,1} = w_{i,1} - \eta \delta_o(\kappa_1) a_i$
 - $v_{i,j} = v_{i,j} - \eta \delta(hidden_j) x_i$


2. Binary classification, take one

- Like Marsland, and regression, for loss use (SE):
$$L_{SE}(\mathbf{t}, \mathbf{y}) = \frac{1}{2} \sum_{j=1}^N (t_j - y_j)^2$$
- The only difference to regression is the logistic activation function: $y = \sigma(x) = \frac{1}{1+e^{-x}}$
- Since the derivative of this is $y(1 - y)$, we get
- $\delta_o(\kappa_1) = (y - t)y(1 - y)$
- The rest is as for regression




2. Binary classification, take two

- Use instead cross-entropy loss (cf. Lecture 7, Marsland 4.6.6)
- $\frac{\partial}{\partial y} L_{CE}(t, y) = -\frac{(t-y)}{y(1-y)}$
- Logistic activation
- $\delta_o(\kappa_1) = -\frac{(t-y)}{y(1-y)} y(1-y) = (y - t)$
- The rest is as for regression and take one




3. Multi-label classification

- Several output nodes
- Logistic activation function
- The model Marsland considers
- $L_{SE}(\mathbf{t}, \mathbf{y}) = \frac{1}{2} \sum_{j=1}^N (t_j - y_j)^2$
 - (The index j here, runs over the output nodes.)
 - We still look at one input only




3. Multi-label classification

- (SE loss, logistic output activation)
- We compute a delta term at each output node, κ_j :
- $\delta_o(\kappa_j) = (y_j - t_j)y_j(1 - y_j)$



3. First layer

- (SE loss, logistic output activation)
- $\delta(hidden_j) =$
- $a_j(1 - a_j) \sum_{i=1}^n \delta_o(\kappa_i) w_{j,i}$
- i.e., sum of delta at output weighted by the connections between them
- The rest as for the others




Putting it together: the Algorithm

- Use the loss function and the derivative of the activation function to compute the delta term at the final node(s),
 - here: $\delta_o(\kappa_j) = (y_j - t_j)y_j(1 - y_j)$ for each node κ_j for $j = 1, \dots, n$
- Compute the delta terms for each node in the hidden layer,
 - here: $\delta(hidden_j) = a_j(1 - a_j) \sum_{i=1}^n \delta_o(\kappa_i)w_{j,i}$ for $j = 1, \dots, k$
- Update the weights by the deltas in both layers
 - $w_{i,j} = w_{i,j} - \eta \delta_o(\kappa_j)a_i$
 - $v_{i,j} = v_{i,j} - \eta \delta(hidden_j)x_i$

By the way:

$$\begin{bmatrix} x_0 & \boxed{x_1} & x_2 & \cdots & x_m \end{bmatrix} \begin{bmatrix} w_{0,1} & w_{0,2} & \cdots & w_{0,n} \\ \boxed{w_{1,1}} & \boxed{w_{1,2}} & \cdots & \boxed{w_{1,n}} \\ w_{2,1} & w_{2,2} & \cdots & w_{2,n} \\ \vdots & \vdots & \ddots & \vdots \\ w_{m,1} & w_{m,2} & \cdots & w_{m,n} \end{bmatrix} = \begin{bmatrix} z_1 & z_2 & \cdots & z_n \end{bmatrix}$$

- To calculate $\sum_{j=1}^m w_{l,j} \delta_j$ by matrices, use
- $[\delta(\kappa_1), \delta(\kappa_2), \dots, \delta(\kappa_n)] \mathbf{W}^T$




Congratulation!

- You just survived backpropagation!
- You now deserve a break and cake!



8.4 Finer details

IN3050/IN4050 Introduction to Artificial Intelligence
and Machine Learning



Practical advices

- Scaling
- Initializing the weights
- Local minima
- Early stopping
- Batch, stochastic, mini-batch
- Number of hidden nodes and hidden layers?
- Activation functions



Scaling

- The $z = \mathbf{w} \cdot \mathbf{x}$ shouldn't be too large for this to work, roughly $|z|$ shouldn't be much more than 1
- For example, normalization (scikit: standardscaler) of each feature

Normalization


- Training data, dimension i :
 $X_i = \{x_{1i}, x_{2i}, \dots, x_{Ni}\}$.
- Let m_i be the corresponding mean value:
 - $m_i = \frac{1}{N} \sum_{j=1}^N x_{ji}$
- Let s_i be the standard deviation
- Define $scale_i(x_{ji}) = \frac{x_{ji} - m_i}{s_i}$
- Use the same scaler on all test data!

Initializing the weights

- The weights:
 - should **not** be initialized to 0
 - should be initialized to random numbers
 - should be initialized to numbers between -1 and 1
- In addition, Marsland recommends to multiply with $\frac{1}{\sqrt{m}}$
 - where m is the number of input nodes


Local minima

- The loss function for MLP is **not convex**
- It can be caught in local minima
- Hence:
 - Make several runs with different initializations and compare the results (mean and std.dev.)
 - Consider methods for escaping local minima, cf. lecture 2 and adding momentum



Early stopping


- The loss on the training data will decrease during training
- There is a danger of overfitting by training for too long:
 - The network knows the training set very well
 - but does not generalize



- Use a validation set V different from the training set.
- After k rounds for some fixed k (e.g., 100):
 - check the loss on V
 - if the loss starts to increase, stop training!

Variations of gradient descent


- Mini-batch training:
 - Pick a subset of the training set of a certain size
 - Calculate the loss for this subset
 - Make one move in the opposite direction of this gradient
- Batch training
 - Use the whole training set in each update
- Stochastic gradient descent:
 - Pick one datapoint at random and update




- SGD/Mini-batch can be a way to avoid local minima

Number of hidden nodes and hidden layers?



- Very much an empirical question
- Use an independent validation set
- Run with different settings and evaluate on the validation set
- Choose the settings which give the best result
- Called **hyper-parameter tuning**
 - (The hyper-parameters are the parameters that you have to set.)



Alternative activation functions in the hidden layer




- There are alternative activation functions
- One may use different functions at different layers
- $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$
- $ReLU(x) = \max(x, 0)$
- ReLU is the preferred method in deep networks



8.5 More on evaluation

IN3050/IN4050 Introduction to Artificial Intelligence
and Machine Learning



Evaluation measures

		Is in C	
		Yes	NO
Class ifier	Yes	tp	fp
	No	fn	tn

- Accuracy: $(tp+tn)/N$
- Precision: $tp/(tp+fp)$
- Recall: $tp/(tp+fn)$

- F-score combines P and R

- $F_1 = \frac{2PR}{P+R} \left(= \frac{1}{\frac{1}{R} + \frac{1}{P}} \right)$

- F_1 called “harmonic mean”

- General form

- $F = \frac{1}{\alpha \frac{1}{P} + (1-\alpha) \frac{1}{R}}$

- for some $0 < \alpha < 1$

Confusion matrix

		<i>gold standard labels</i>		
		gold positive	gold negative	
<i>system output labels</i>	system positive	true positive	false positive	precision = $\frac{tp}{tp+fp}$
	system negative	false negative	true negative	
		recall = $\frac{tp}{tp+fn}$		accuracy = $\frac{tp+tn}{tp+fp+tn+fn}$

Figure 6.4 Contingency table

- Beware what the rows and columns are:
 - Marsland swaps them

Confusion matrix

		<i>gold labels</i>			
		urgent	normal	spam	
<i>system output</i>	urgent	8	10	1	$\text{precision}_u = \frac{8}{8+10+1}$
	normal	5	60	50	$\text{precision}_n = \frac{60}{5+60+50}$
	spam	3	30	200	$\text{precision}_s = \frac{200}{3+30+200}$
		$\text{recall}_u = \frac{8}{8+5+3}$	$\text{recall}_n = \frac{60}{10+60+30}$	$\text{recall}_s = \frac{200}{1+50+200}$	

- Precision, recall and f-score can be calculated for each class against the rest

Figure 6.5 Confusion matrix for a three-class categorization task, showing for each pair of classes (c_1, c_2), how many documents from c_1 were (in)correctly assigned to c_2


3. Multi-[label | class: ovr] classification

Task:

- A set of labels $\mathcal{C} = \{C_1, C_2, \dots, C_n\}$
- Multi-class, $f(\mathbf{x}) = C_i$ for an $C_i \in \mathcal{C}$
- Multi-label, $f(\mathbf{x})(C_i) = 1$ or 0
for each $C_i \in \mathcal{C}$

Representations:

- Multi-class, e.g., $(0,0,1,0, \dots 0)$
- Multi-label, e.g., $(1,0,1,0, \dots 1)$



3. Multi-[label | class: ovr] classification

Task:

- A set of labels $\mathcal{C} = \{C_1, C_2, \dots, C_n\}$
- Multi-class, $f(\mathbf{x}) = C_i$ for an $C_i \in \mathcal{C}$
- Multi-label, $f(\mathbf{x})(C_i) = 1$ or 0
for each $C_i \in \mathcal{C}$

Representations:

- Multi-class, e.g., $(0,0,1,0, \dots 0)$
- Multi-label, e.g., $(1,0,1,0, \dots 1)$

• Model (common):

- $y_j = \frac{1}{1+e^{-z_j}}$
- MSE-loss

• Same learning by backpropagation

• Difference in application:

- Multi-class, one-vs-rest: $f(\mathbf{x}) = C_i$,
where $i = \operatorname{argmax}_{j=1,\dots,n} y_j$
- Multi-label: $f(\mathbf{x})(C_i) = y_i > 0.5$

4. Multinomial Logistic Regression

- Also called **softmax-classifier**

- $y_j = \frac{e^{z_j}}{\sum_{k=1}^n e^{z_k}}$

- Cross-entropy loss:

- $L_{CE}(\mathbf{y}, \mathbf{t}) = -\sum_{j=1}^n t_j \log y_j$


- Remember:

- $\frac{\partial}{\partial w_{i,j}} L_{CE}(\mathbf{x}, \mathbf{t}, \mathbf{w}) = (y_j - t_j) a_i$

- Similarly:


- $\frac{\partial}{\partial a_i} L_{CE}(\mathbf{x}, \mathbf{t}, \mathbf{w}) = (y_j - t_j) w_{i,j}$

- $\delta_o(\kappa_j) = (y_j - t_j)$



4. Multinomial Logistic Regression

- $\delta_o(\kappa_j) = (y_j - t_j)$
- As before:
 - $\delta(hidden_j) =$
 - $a_i(1 - a_i) \sum_{j=1}^n \delta_o(\kappa_j) w_{i,j}$
- Update the weights:
 - $w_{ij} = w_{i,j} - \eta \delta_o(\kappa_j) a_{i,j}$
 - $v_{i,j} = v_{i,j} - \eta \delta(hidden_j) x_i$
- Remark the modularity:
 - Composite functions
 - Partial derivatives



The importance of Multinomial Log.Reg

- The multinomial logistic regression, or **softmax classifier** is an essential tool in modern (deep) neural networks.
- E.g., **Natural Language Processing**:
 - Language modelling: Which word comes next? :
 - *I like to eat ...*
 - Softmax over all English words
 - Translation
 - Translate *back* into Norwegian
 - Softmax over candidates: {*bak, rygg, back, støtt, rygge, støtte,*}
 - Tagging
 - What is the part of speech for *back* in *They back the proposition*
 - Softmax over {**Noun, Verb, Adj., Preposition, ...**}